

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn import metrics
```

```
In [2]: car_dataset = pd.read_csv(r"C:\Users\robin\Downloads\car_data.csv")
```

```
In [3]: car_dataset.head(3)
```

```
Out[3]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Trans
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	

```
In [4]: # cheaking number of row and columns
car_dataset.shape
```

```
Out[4]: (301, 9)
```

```
In [5]: car_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Car_Name        301 non-null   object
1   Year            301 non-null   int64
2   Selling_Price   301 non-null   float64
3   Present_Price   301 non-null   float64
4   Kms_Driven      301 non-null   int64
5   Fuel_Type       301 non-null   object
6   Seller_Type     301 non-null   object
7   Transmission    301 non-null   object
8   Owner           301 non-null   int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

```
In [6]: # checking the number of missing values
car_dataset.isnull().sum()
```

```
Out[6]: Car_Name      0
        Year          0
        Selling_Price  0
        Present_Price  0
        Kms_Driven     0
        Fuel_Type      0
        Seller_Type     0
        Transmission   0
        Owner          0
        dtype: int64
```

```
In [8]: # number of values for categorical values
print(car_dataset.Fuel_Type.value_counts())
print(car_dataset.Seller_Type.value_counts())
print(car_dataset.Transmission.value_counts())
```

```
Petrol    239
Diesel     60
CNG        2
Name: Fuel_Type, dtype: int64
Dealer     195
Individual 106
Name: Seller_Type, dtype: int64
Manual     261
Automatic   40
Name: Transmission, dtype: int64
```

```
In [9]: #encoding the categorical data
# we have label encoding categorical value replace the 1 and 0
car_dataset.replace({'Fuel_Type':{'Petrol':0, 'Diesel':1, 'CNG':2}},inplace=True)
car_dataset.replace({'Seller_Type':{'Dealer':0, 'Individual':1}},inplace=True)
car_dataset.replace({'Transmission':{'Manual':0, 'Automatic':1}},inplace=True)
```

```
In [11]: car_dataset.head(3)
```

```
Out[11]:
```

	Car_Name	Year	Selling_Price	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmission
0	ritz	2014	3.35	5.59	27000	0	0	0
1	sx4	2013	4.75	9.54	43000	1	0	0
2	ciaz	2017	7.25	9.85	6900	0	0	0

```
In [12]: #splitting our original data into train and test
x = car_dataset.drop(columns=['Car_Name', 'Selling_Price'],axis=1)
y = car_dataset['Selling_Price']
```

In [14]: `print(x)`

	Year	Present_Price	Kms_Driven	Fuel_Type	Seller_Type	Transmissio
0	2014	5.59	27000	0	0	
1	2013	9.54	43000	1	0	
2	2017	9.85	6900	0	0	
3	2011	4.15	5200	0	0	
4	2014	6.87	42450	1	0	
..	
296	2016	11.60	33988	1	0	
297	2015	5.90	60000	0	0	
298	2009	11.00	87934	0	0	
299	2017	12.50	9000	1	0	
300	2016	5.90	5464	0	0	

	Owner
0	0
1	0
2	0
3	0
4	0
..	...
296	0
297	0
298	0
299	0
300	0

[301 rows x 7 columns]

In [15]: `print(y)`

0	3.35
1	4.75
2	7.25
3	2.85
4	4.60
...	
296	9.50
297	4.00
298	3.35
299	11.50
300	5.30

Name: Selling_Price, Length: 301, dtype: float64

```
In [16]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.1,random_state=42)
# 0.1 means 10% in test data and 90% in train
```

```
In [17]: print(x.shape,x_train.shape,x_test.shape)
```

```
(301, 7) (270, 7) (31, 7)
```

```
In [18]: print(y.shape,y_train.shape,y_test.shape)
```

```
(301,) (270,) (31,)
```

Linear Regression

```
In [19]: linear_regression = LinearRegression()
```

```
In [20]: #training the linear_regression machine model
linear_regression.fit(x_train,y_train)
```

```
Out[20]: LinearRegression()
```

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.

On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [21]: # prediction on training data
x_train_prediction = linear_regression.predict(x_train)
```


```
In [22]: # Assuming y_train and x_train_prediction are defined
error_score = metrics.r2_score(y_train, x_train_prediction)
```

```
In [23]: print("R squared Error:",error_score)
```

```
R squared Error: 0.8799451660493701
```

```
In [24]: ### Visualize the actual price and predicted price
```

```
In [25]: plt.scatter(y_train,x_train_predicton)
plt.xlabel("Actual price ")
plt.ylabel("predicted price")
plt.title("actual prices vs predicted prices")
plt.show()
```




```
In [26]: ## Evaluate the test model
```

```
In [27]: # prediction on test data
test_data_predicton = linear_regression.predict(x_test)
```

```
In [28]: # Assuming y_train and x_train_prediction are defined
error_score = metrics.r2_score(y_test, test_data_predicton)
print("R squared Error:",error_score)
```

R squared Error: 0.8365766715026928

```
In [29]: plt.scatter(y_test, test_data_prediction)
plt.xlabel("Actual price ")
plt.ylabel("predicted price")
plt.title("actual prices vs predicted prices")
plt.show()
```



```
In [30]: input_data = (2017,9.85,6900,0,0,0,0)

#changing the input_data type list to numpy array
input_data_as_numpy_array = np.array(input_data)

#reshape the np as we are predicting for one instance
input_data_reshape= input_data_as_numpy_array.reshape(1,-1)

prediction = linear_regression.predict(input_data_reshape)

print(prediction)
```

[7.07156998]

C:\Users\robin\anaconda3\lib\site-packages\sklearn\base.py:465: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
warnings.warn(

```
In [31]: ## Lasso Regression
lass_reg_model = Lasso()
```

```
In [32]: #training the linear_regression machine model
lass_reg_model.fit(x_train,y_train)
```

Out[32]: Lasso()

In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.


On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.

```
In [33]: Traning_data_prediction =lass_reg_model.predict(x_train)
```

```
In [34]: # Assuming y_train and x_train_prediction are defined
error_score1 = metrics.r2_score(y_train, Traning_data_prediction)
print("R squared Error:",error_score1)
```

R squared Error: 0.8427856123435794

```
In [35]: plt.scatter(y_train,Traning_data_prediction)
plt.xlabel("Actual price ")
plt.ylabel("predicted price")
plt.title("actual prices vs predicted prices")
plt.show()
```




```
In [36]: # prediction on test data
test_data_predicton = lass_reg_model.predict(x_test)## Evaluate the test r
```

```
In [37]: # Assuming y_train and x_train_prediction are defined
error_score = metrics.r2_score(y_test, test_data_predicton)
print("R squared Error:",error_score)
```

R squared Error: 0.8709167941173195

```
In [38]: plt.scatter(y_test,test_data_predicton)
plt.xlabel("Actual price ")
plt.ylabel("predicted price")
plt.title("actual prices vs predicted prices")
plt.show()
```




```
In [39]: input_data = (2017,9.85,6900,0,0,0,0)

#changing the input_data type list to numpy array
input_data_as_numpy_array = np.array(input_data)

#reshape the np as we are predicting for one instance
input_data_reshape= input_data_as_numpy_array.reshape(1,-1)

prediction = lass_reg_model.predict(input_data_reshape)

print(prediction)
```

```
[7.16105963]
```

```
C:\Users\robin\anaconda3\lib\site-packages\sklearn\base.py:465: UserWarning: X does not have valid feature names, but Lasso was fitted with feature names
  warnings.warn(
```

```
In [ ]:
```